

Midterm Review

CMPUT 365

Fall 2021

Nov 15

- Hopefully you had a relaxing Reading Week!
- Midterm this Friday, here in ETLC with pen and paper for most:
 - A few students will take the midterm over Zoom and E-class

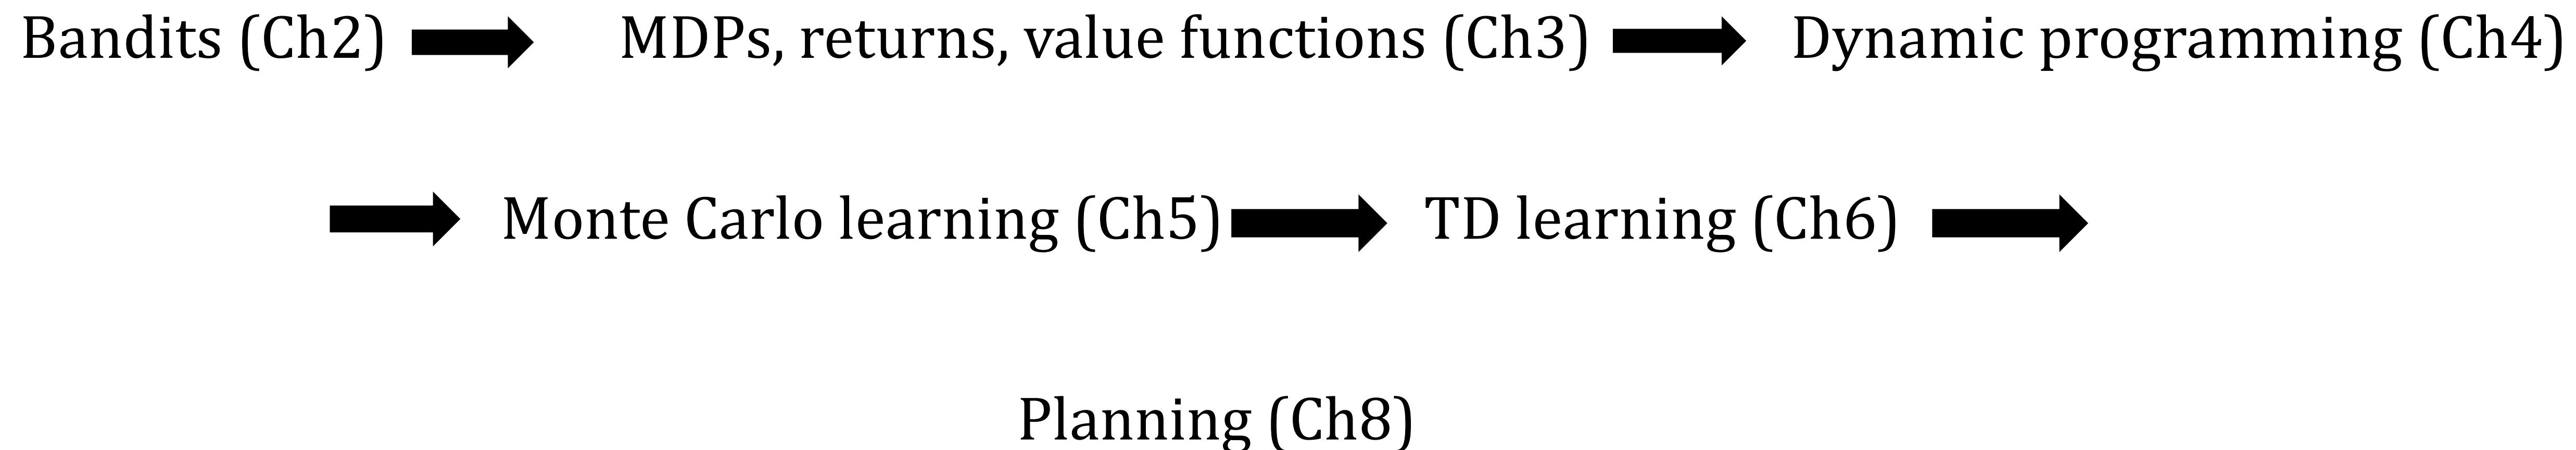
Midterm details

- **If you are taking the test here in ETLC**
 - you can prepare a one-page double sided cheat-sheet
 - you can use an electronic device to access a pdf of the textbook
 - no other files (e.g., no course slides, worksheets or videos)
 - no accessing the internet in anyway
 - exam will be on paper (bring a pen / pencil)
 - physically raise hand if you have questions

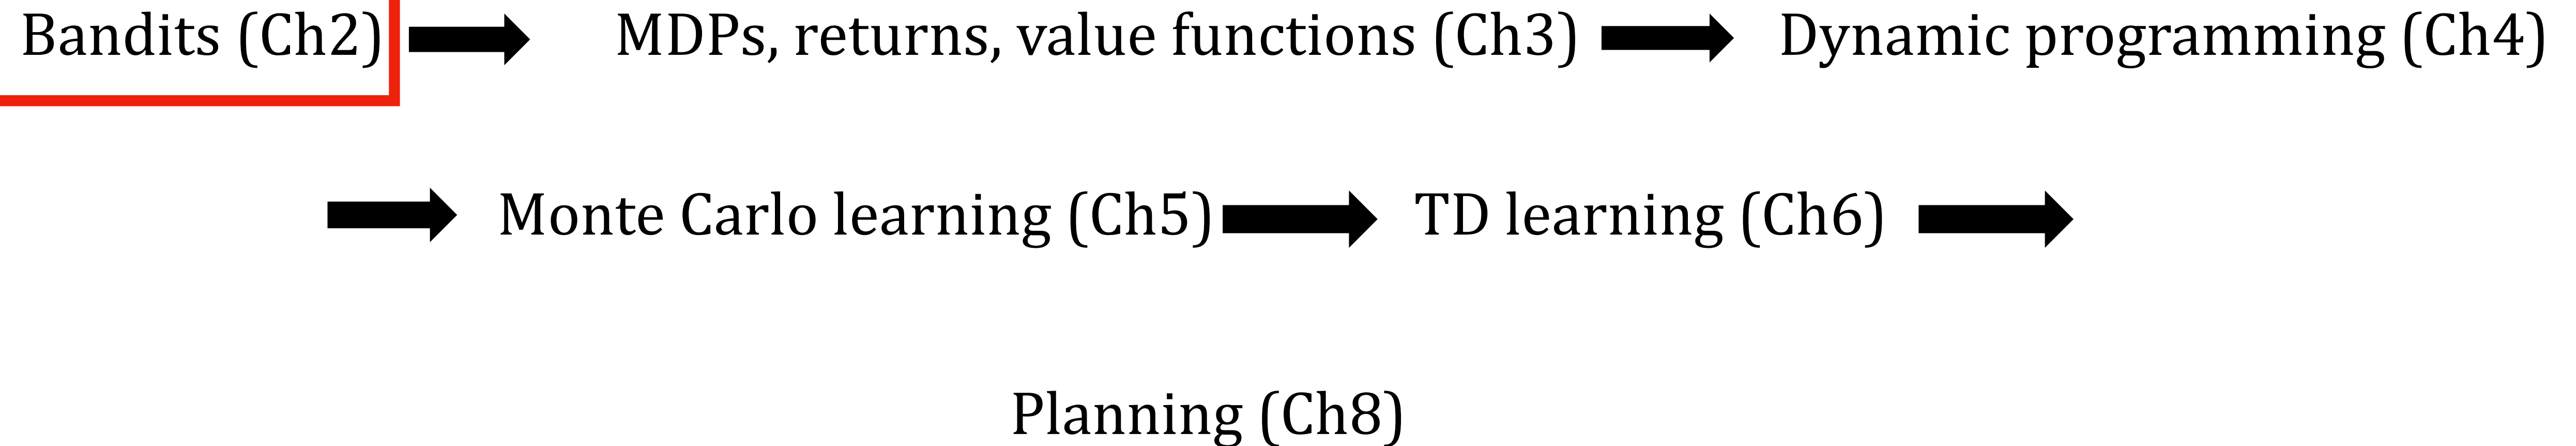
Midterm details

- **If you are taking the test remotely**
 - same as in-person but in addition ...
 - exam will be via e-class
 - **must join regular friday zoom call and have your camera on showing your face at all times during the test** (turning off the camera will be considered cheating)
 - NO COMMUNICATING IN ANY WAY WITH OTHER STUDENTS
 - TAs will monitor the zoom call (and it will be recorded)
 - use private zoom message to one of the TAs if you have questions

Course Roadmap




Course Roadmap



Bandits

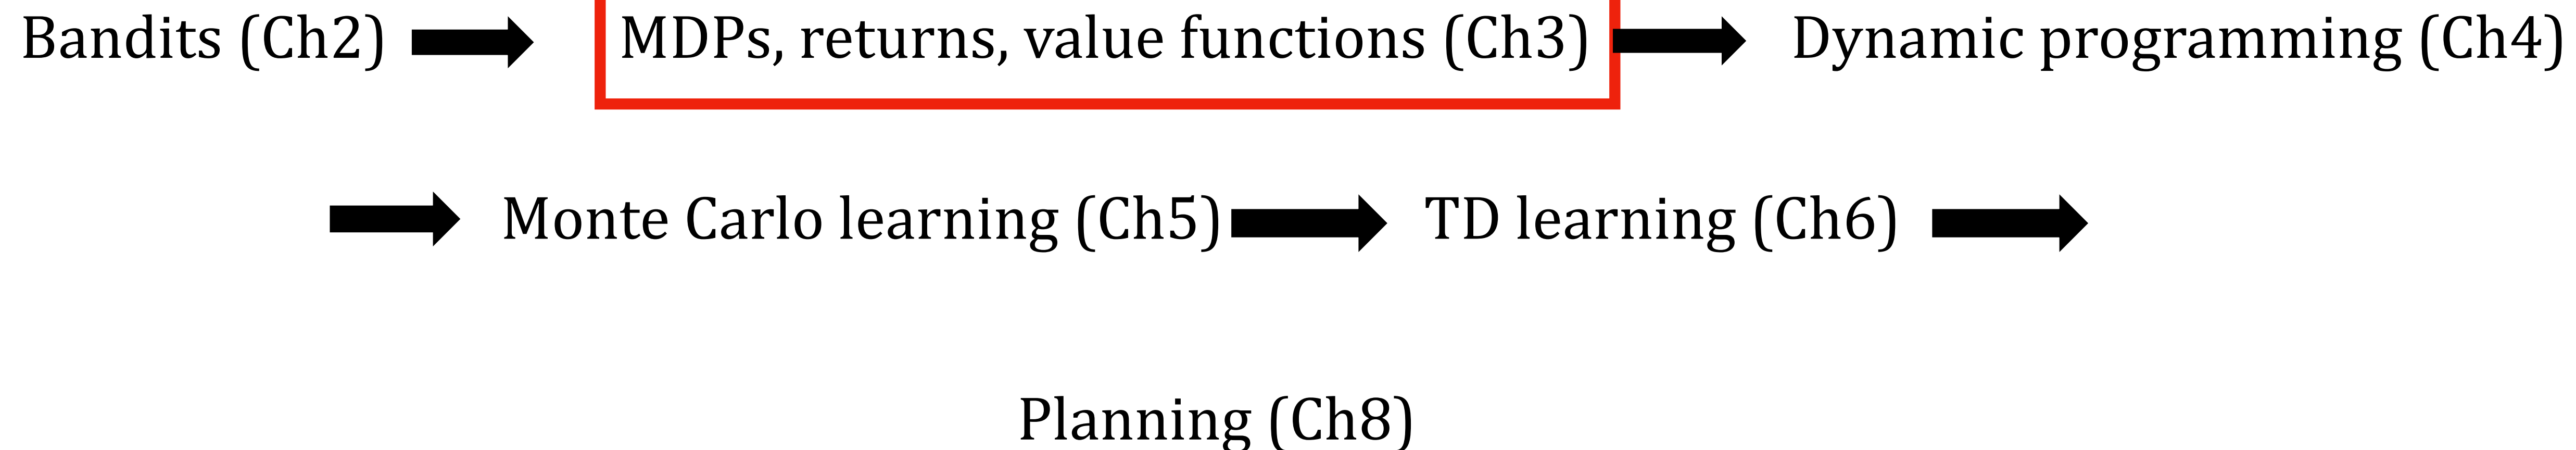
- Simple decision making problem with 1 state



Bandits


- Know the exploration-exploitation tradeoff!
 - i.e. Why shouldn't you always be greedy? Why not constantly explore?
- Know about incremental averaging (and why we do it!)
 - $\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$

Course Roadmap



MDPs, Returns, Value Functions

- Decision making problems with many states



MDPs, Returns, Value Functions

- Sequential decision making: must take many actions in a row to maximize reward
- Agent is concerned with **returns**:

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{k-1} R_{t+k} + \dots$$

- Specifically, the agent estimates the **expected return**, which depends on the agent's **policy** and the **environment dynamics**

MDPs, Returns, Value Functions

- Value-based methods address this by learning to predict the **expected return**, i.e. learning **value functions**
- Value functions:

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi} [G_t \mid S_t = s] \quad \text{“How good is this state”}$$

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi} [G_t \mid S_t = s, A_t = a] \quad \text{“How good is taking **this action** in this state”}$$

MDPs, Returns, Value Functions

- Bellman Equations: write the value of a state in terms of the value of another state
- i.e. for all states:

$$\begin{aligned} v_{\pi}(s) &\doteq \mathbb{E}_{\pi} [G_t \mid S_t = s] \\ &= \sum_a \pi(a \mid s) \sum_{s'} \sum_r p(s', r \mid s, a) [r + \gamma v_{\pi}(s')] \end{aligned}$$

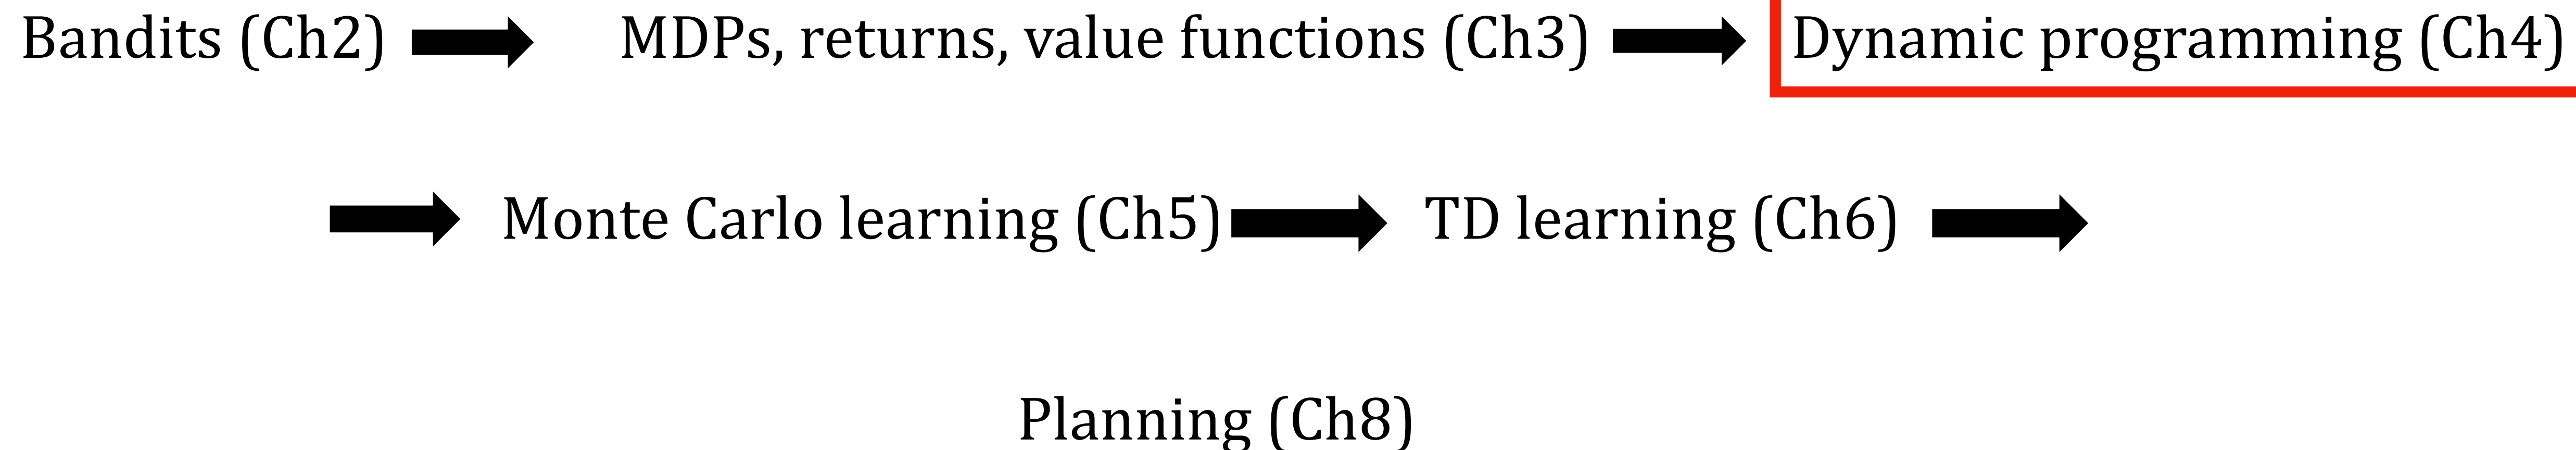
MDPs, Returns, Value Functions

- **Policy improvement**
 - If you derive a **greedy** policy with respect to the *action-values* of another policy, the new policy will be at least as “good” as the previous one
 - If the new policy did not change from the previous policy, the policy is greedy with respect to its **own** value function, and is an optimal policy π^*
 - Optimal value functions denoted $v^*(s)$ and $q^*(s,a)$

Self-test: Connections to Course 2

- We only talked about the policy improvement result in Course 1, when we did DP
- How is policy improvement relevant for the sample-based methods in Course 2?
 - how is it relevant for Sarsa?
 - how is it relevant for Q-learning?

Course Roadmap



Dynamic Programming: Iterative Policy Evaluation

- Computes an approximate value function $V(s) \approx v_{\pi}(s)$
- **Sweeps** across all states and actions, and evaluates the Bellman equation using the current estimates in the value function

$$v_{k+1}(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_k(s')]$$

Dynamic Programming

- Introduces the idea of **bootstrapping** - basing the update to a state's value on the agent's *current value estimates* of successor states
- Requires knowledge of the **environment dynamics** $p(s', r \mid s, a)$
 - this is a model-based method since it assume access to the environment model p

Dynamic Programming: Value Iteration

- Uses Bellman equation to iterate towards v^* (and so towards π^*)


- $$v_{k+1}(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma v_k(s')]$$

- Contrast with policy evaluation update

- $$v_{k+1}(s) \leftarrow \sum_a \pi(a | s) \sum_{s',r} p(s', r | s, a) [r + \gamma v_k(s')]$$

- Policy iteration uses greedy policy, and fully evaluates the values for that policy (multiple sweeps to do IPE)
- Value Iteration greedifies, after only one sweep of evaluating with the current greedy policy

Course Roadmap



Monte Carlo Learning

- Policy Evaluation: estimates the value function $V(s) \approx v_{\pi}(s)$
- Sample *returns* from states by **following policy π** , then average those *returns* for each state


Monte Carlo Learning

- Doesn't need a model of the environment
- We only used it in episodic problems: learning only occurs **after** each episode

Monte Carlo Self-test


- We talked about two versions of Monte Carlo for prediction. The first uses a sample average (sample mean) of returns from a state s . The second uses the following incremental update rule, for a constant stepsize $\alpha > 0$
- $V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$
- What is the primary difference between the values learned with the sample average and those with the incremental update rule?

Course Roadmap




Temporal Difference Learning

Previous experience:




Temporal Difference Learning

Previous experience:




Temporal Difference Learning

New experience:



Temporal Difference Learning

New experience:



TD Learning

- Estimates value function $V(s) \approx v_{\pi}(s)$
- TD is for Policy Evaluation (prediction), Sarsa and Q-learning are TD variants for control
- Combines ideas from Monte Carlo and Dynamic Programming - uses a mix of sampled information and bootstrapping off of current estimates
- Can learn *online*, without having to wait for the end of an episode

TD Learning Algorithms

One-step TD (or TD(0)):

$$V(S_t) \leftarrow V(S_t) + \alpha [\hat{G}_t - V(S_t)]$$

$$\hat{G}_t = R_{t+1} + \gamma V(S_{t+1})$$

One-step Sarsa (or Sarsa(0)):

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [\hat{G}_t - Q(S_t, A_t)]$$

$$\hat{G}_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$$

Self-test: What makes Sarsa a control algorithm?

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Can you get a policy evaluation variant of Sarsa? (i.e., TD that estimates action-values)

We would call this Sarsa for Prediction. By default, Sarsa means Sarsa for Control

TD Self-test

- What is the difference between **online** and **offline** updating?
- What is the difference between Sarsa, Q-learning, and Expected Sarsa?

Self-test: Variance in Updates

- We talked a bit about the variance in updates
- Why does variance in the update matter?
- Which do you think will have a lower variance update: MC or TD, for prediction?
- Which do you think will have a lower variance update: Sarsa or Expected Sarsa?

Course Roadmap

Bandits (Ch2) ➡ MDPs, returns, value functions (Ch3) ➡ Dynamic programming (Ch4)
➡ Monte Carlo learning (Ch5) ➡ TD learning (Ch6) ➡

Planning (Ch8)

Planning, Learning and Acting

- Planning: a process which takes a **model** as input and produces or improves a **policy**
- Dyna uses a model to **simulate experience** and improve its value estimates, where greedifying with respect to these value estimates produces an improved **policy**

.

Dyna Self-test


- What is a model? (where model has a technical definition for RL and this course)
 - A procedure that produces a possible next state and reward, for a given state and action
- What is the difference between **simulated** and **real** experience?
- Why can a distribution model be used as a sample model? How do we sample?

Course Roadmap

Bandits (C2) ➡ MDPs, returns, value functions (C3) ➡ Dynamic programming (C4)

➡ Monte Carlo learning (C5) ➡ TD learning (C6) ➡

Planning (C8) ➡ Function approximation (C9)



Comments about Test Answers

- The goal is to see your thought process. Try to explain your answer clearly and give reasons, rather than just writing down an answer (but still be concise)
- Don't vomit on the page: if you write many answers, and some of them are wrong, I will mark the wrong ones
- I don't give partial marks for wrong answers, but if you demonstrate understanding (but its wrong) then I might give partial marks for that
- **One common problem:** answering a different question than is asked
 - after answering, re-read the question and check if you answered it
 - If you think: "this is a simple question, he must have meant this other more complicated thing", you are probably wrong. I probably meant the simple thing.